

# How to handle lots of requests and not get crazy along the way

- Damian Kampik, u2i

# Short intro

- 5 yrs of experience as Ruby programmer
- 1,5 yrs at u2i
- We're just around the corner - come visit us!



u2i



# QR Codes



# AGENDA

---

# Agenda

1. What's the problem?
2. How did we solve it?
  - a. The idea
  - b. Lambda
  - c. Sidekiq
  - d. Apex
3. Summary

# PROBLEM

---

What did we need to implement?

# What did we need to implement?

- SMS emergency alerts




# What did we need to implement?




- SMS emergency alerts
  - Maximize the number of successfully received messages

# What did we need to implement?

- SMS emergency alerts
  - Maximize the number of successfully received messages
  - Collect delivery statistics




Twilio








Twilio Notify



Twilio



# Twilio Notify

```
def service
```

```
  return @service unless @service.nil?
```

```
  account_sid = ENV['TWILIO_ACCOUNT_SID']
```

```
  auth_token = ENV['TWILIO_AUTH_TOKEN']
```

```
  client = Twilio::REST::Client.new(account_sid, auth_token)
```

```
  @service = client.notify.v1
```

```
    .services(ENV['TWILIO_NOTIFY_SERVICE'])
```

```
end
```

Notify  
docs



# Twilio Notify

```
def phones_binding
  @phones.map do |phone|
    {
      binding_type: :sms,
      address: phone
    }.to_json
  end
end
```

Notify  
docs



# Twilio Notify


```
def send_message
  service.notifications.create({
    to_binding: phones_binding,
    body: @content
  })
end
```

Notify  
docs









Twilio Notify





Twilio
















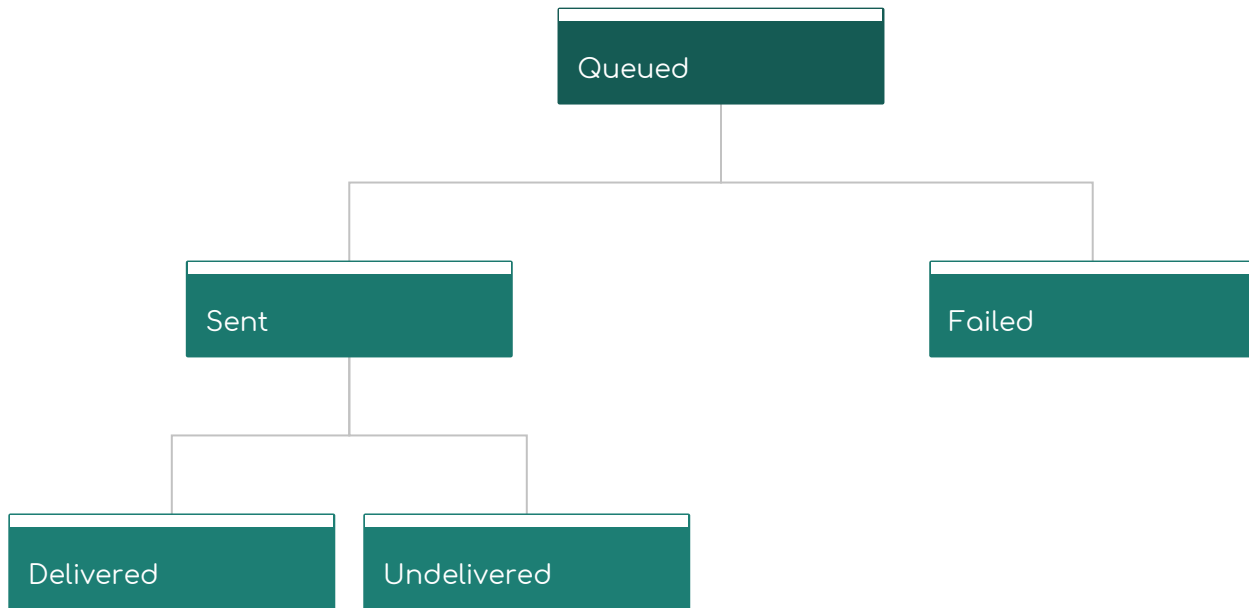




Twilio




# Diagram of message statuses










Status  
callbacks  
docs







Twilio









Twilio




Status x3









Status x3






Twilio









+ 2997 remaining recipients












Twilio



+ 2997 remaining recipients




Twilio



+ 2997 remaining recipients




9k requests




Twilio


Status x3



Status x3




Status x3



+ 2997 remaining recipients




9k requests  
DoS




Twilio


Status x3




Status x3



Status x3




+ 2997 remaining recipients






9k requests  
DoS




Twilio


Status x3



Status x3



Status x3



+ 2997 remaining recipients

Thx, Twilio :(

What was the problem?




# What was the problem?

- Too many requests from Twilio interfering with regular website traffic

# What was the problem?

- Too many requests from Twilio interfering with regular website traffic
- We could not display correct delivery statistics

## Message delivery rate




# SOLUTION

---

# AWS Lambda





# Why Lambda?

# Why Lambda?

- Load on AWS side



# Why Lambda?

- Load on AWS side
- High scalability - up to 1000 simultaneous executions

# Why Lambda?

- Load on AWS side
- High scalability - up to 1000 simultaneous executions
- We're paying for the infrastructure only when we have to handle incoming Twilio requests

# Why Lambda?

- Load on AWS side
- High scalability - up to 1000 simultaneous executions
- We're paying for the infrastructure only when we have to handle incoming Twilio requests
- The feature could be written in a few simple methods

# Why Sidekiq?

## Why Sidekiq?

- To enqueue status update data, you only have to create a record with job metadata in Redis

## Why Sidekiq?

- To enqueue status update data, you only have to create a record with job metadata in Redis
- Servers will be handling status updates when it has free resources

## Why Sidekiq?

- To enqueue status update data, you only have to create a record with job metadata in Redis
- Servers will be handling status updates when it has free resources
- You can assign workers to process message status updates exclusively

## Why Sidekiq?

- To enqueue status update data, you only have to create a record with job metadata in Redis
- Servers will be handling status updates when it has free resources
- You can assign workers to process message status updates exclusively

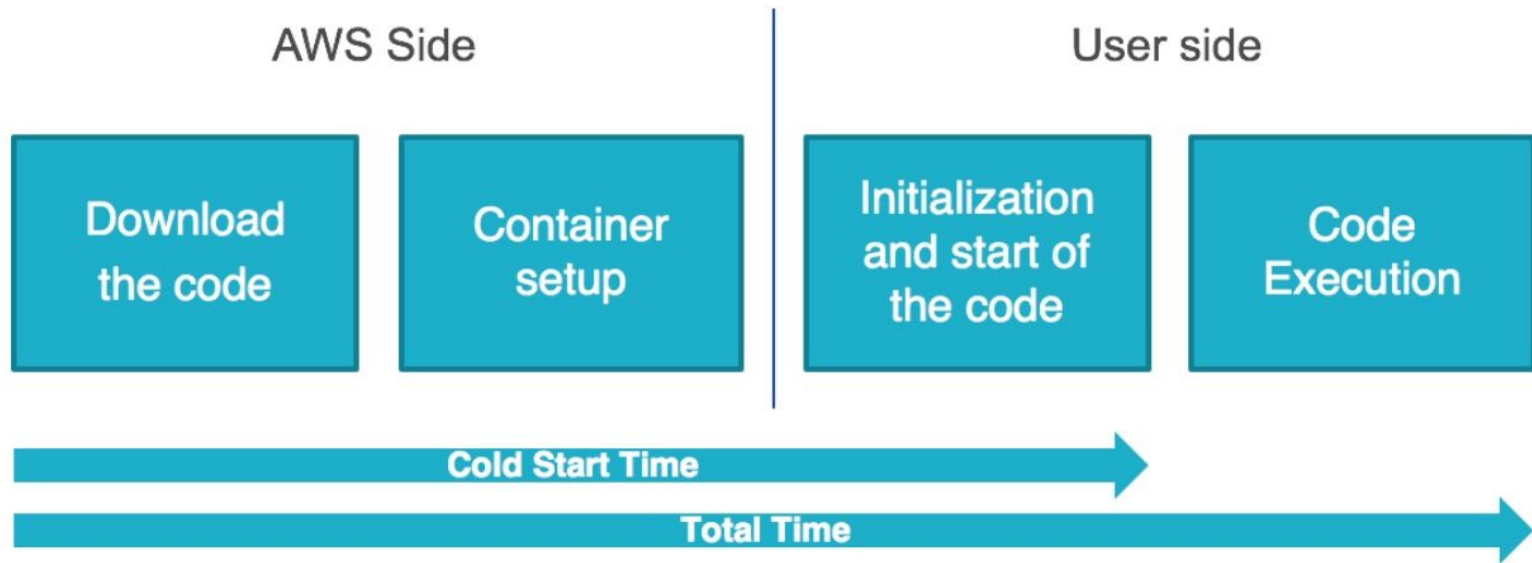


What have we learnt the hard way?

What have we learnt the hard way?



# Cold start - behind the scenes



Source : AWS re:Invent 2017 – Become a serverless Black Belt

How to bypass the issue?

## How to bypass the issue?

- Declare lowest amount of RAM possible

## How to bypass the issue?

- Declare lowest amount of RAM possible
- Use interpreted language

# How to bypass the issue?

- Declare lowest amount of RAM possible
- Use interpreted language
- Optimize memory consumption of executed code

# How to bypass the issue?

- Declare lowest amount of RAM possible
- Use interpreted language
- Optimize memory consumption of executed code
- Optimize size of the code



# IMPLEMENTATION

---

# Twilio request validation

```
def handle(event, context):  
    if 'headers' in event and 'X-Twilio-Signature' in  
        event['headers'] and 'body' in event:  
        twilio_signature =  
            event['headers']['X-Twilio-Signature']  
        form_parameters = dict(parse_qs(event['body']))  
  
        if request_valid(form_parameters, twilio_signature)  
            and u'From' in form_parameters:  
  
    ...
```

# Twilio request validation

```
def request_valid(parameters, signature):  
    validator =  
        RequestValidator(os.environ['TWILIO_AUTH_TOKEN'])  
  
    return validator.validate(  
        os.environ['REQUEST_URL'],  
        parameters,  
        signature  
    )
```

# Sidekiq thru Python

```
def push_to_sidekiq(parameters):  
    elasticache = redis.StrictRedis(  
        host=os.environ['ELASTICACHE_URL'], port=6379, db=0)  
  
    args = [{  
        "job_class": "StatusCallbackJob",  
        "job_id": str(uuid.uuid4()),  
        "queue_name": "callbacks",  
        "locale": "en",  
        "arguments": [parameters['To'], parameters['MessageStatus'],  
                      parameters['Body'], int(time.time())]  
    }]
```

Sidekiq  
FAQ



# Sidekiq thru Python

```
sidekiq_params = {  
    "class": "ActiveJob::QueueAdapters::SidekiqAdapter::JobWrapper",  
    "wrapped": "StatusCallbackJob",  
    "jid": binascii.b2a_hex(os.urandom(12)),  
    "args": args,  
    "created_at": int(time.time()),  
    "enqueued_at": int(time.time()),  
    "queue": "callbacks",  
    "retry": True  
}  
elasticache.lpush("{0}:queue:callbacks".format(  
    os.environ['REDIS_NAMESPACE']), json.dumps(sidekiq_params))
```

Sidekiq  
FAQ




# Worker collecting the data

```
class StatusCallbackJob < ActiveJob::Base
  queue_as :callbacks

  VALID_ALERT_THRESHOLD = 6.hours
  TWILIO_STATUSES = %w[queued failed sent delivered
undelivered].freeze

  def perform(phone, status, content, timestamp)
    ...
  end
end
```

# A quick note about Terraform



HashiCorp

# Terraform


Lambda in  
Terraform  
tutorial



# Apex - efficient Lambda management



Apex  
website





# Project structure

```
.
├── env.json
├── functions
│   ├── alerts
│   │   ├── <libraries>
│   │   ├── main.py
│   │   ├── requirements.txt
│   │   └── setup.cfg
├── <terraform related files>
└── project.json
```

# Project structure

```
.
├── env.json
├── functions
│   ├── alerts
│   │   ├── <libraries>
│   │   ├── main.py
│   │   ├── requirements.txt
│   │   └── setup.cfg
├── <terraform related files>
└── project.json
```

# Project structure

```
.
├── env.json
├── functions
│   ├── alerts
│   │   ├── <libraries>
│   │   ├── main.py
│   │   ├── requirements.txt
│   │   └── setup.cfg
├── <terraform related files>
└── project.json
```

# Project structure

```
.
├── env.json
├── functions
│   ├── alerts
│   │   ├── <libraries>
│   │   ├── main.py
│   │   ├── requirements.txt
│   │   └── setup.cfg
├── <terraform related files>
└── project.json
```

# project.json

```
{  
  "name": "alerts",  
  "nameTemplate": "{{ .Project.Environment }}_{{  
.Function.Name }}",  
  "description": "",  
  "memory": 128,  
  "timeout": 60,  
  "role": <your IAM role>  
  ...  
}
```

# project.json

```
"environment": {  
  "REQUEST_URL": <API Gateway URL to invoke the function>,  
  "ELASTICACHE_URL": <Redis instance URL>,  
  "REDIS_NAMESPACE": <Redis namespace declared in Sidekiq  
config>  
},  
"vpc": {  
  "securityGroups": ["sg-1234567890abcdef"],  
  "subnets": ["subnet-12345678", "subnet-90abcdef"]  
}  
...
```

# project.json

```
"hooks": {  
  "build": "pip install -r requirements.txt -t ."  
}  
}
```

# How to use Apex

```
$ apex deploy --env=<target_env> --env-file=<env_file>  
<function_name>
```



# How to use Apex

```
$ apex deploy --env=<target_env> --env-file=<env_file>  
<function_name>
```

```
$ apex invoke --env=<target_env> --env-file=<env_file>  
<function_name> < event.json
```



LAMBDA


SIDEKIQ

TWILIO

RAILS

APEX

# Message delivery rate



# TAKEAWAYS

---

What to remember?

# What to remember?

- The architecture does not scale indefinitely

# What to remember?

- The architecture does not scale indefinitely
- You can fall victim to cold start

# What to remember?

- The architecture does not scale indefinitely
- You can fall victim to cold start
- This solution will work only for scenarios in which you want to throttle the traffic



# What to remember?

- The architecture does not scale indefinitely
- You can fall victim to cold start
- This solution will work only for scenarios in which you want to throttle the traffic
- From Nov 2018 you can use Ruby on Lambda!

# QUESTIONS?

---

Thanks a lot! :)

---